

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import os
import PIL
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
import pathlib
import pandas as pd
from PIL import Image
from PIL.ImageDraw import Draw
```

```
In [2]: width = 480
height = 480

num_classes = 1
classes = ["sign"]
```

Processing Images and Annotations

```
In [3]: TRAINING_CSV_FILE = 'Data/training_data.csv'
VALIDATION_CSV_FILE = 'Data/validation_data.csv'
TRAINING_IMAGE_DIR = 'Images/training'
VALIDATION_IMAGE_DIR = 'Images/validation'

training_image_records = pd.read_csv(TRAINING_CSV_FILE)
validation_image_records = pd.read_csv(VALIDATION_CSV_FILE)

train_image_path = os.path.join(os.getcwd(), TRAINING_IMAGE_DIR)
validation_image_path = os.path.join(os.getcwd(), VALIDATION_IMAGE_DIR)

train_images = []
train_targets = []
train_labels = []

for index, row in training_image_records.iterrows():

    (filename, width, height, class_name, xmin, ymin, xmax, ymax) = row

    train_image_fullpath = os.path.join(train_image_path, filename)
    train_img = keras.preprocessing.image.load_img(train_image_fullpath, target_size=(width, height))
    train_img_arr = keras.preprocessing.image.img_to_array(train_img)

    # Here we make sure the bounding boxes' vertices are always at the same
    # resolution as the image
    xmin = round(xmin/ width, 2)
    ymin = round(ymin/ height, 2)
    xmax = round(xmax/ width, 2)
    ymax = round(ymax/ height, 2)
```

```

train_images.append(train_img_arr)
train_targets.append((xmin, ymin, xmax, ymax))
train_labels.append(classes.index(class_name))

validation_images = []
validation_targets = []
validation_labels = []

for index, row in validation_image_records.iterrows():

    (filename, width, height, class_name, xmin, ymin, xmax, ymax) = row

    validation_image_fullpath = os.path.join(validation_image_path, filename)
    validation_img = keras.preprocessing.image.load_img(validation_image_fullpath)
    validation_img_arr = keras.preprocessing.image.img_to_array(validation_img)

    xmin = round(xmin/ width, 2)
    ymin = round(ymin/ height, 2)
    xmax = round(xmax/ width, 2)
    ymax = round(ymax/ height, 2)

    validation_images.append(validation_img_arr)
    validation_targets.append((xmin, ymin, xmax, ymax))
    validation_labels.append(classes.index(class_name))

```

Building the Model

In [4]:

```
#create the common input layer
input_shape = (height, width, 3)
input_layer = tf.keras.layers.Input(input_shape)
```

In [5]:

```
base_layers = layers.experimental.preprocessing.Rescaling(1./255, name='bl_1')
base_layers = layers.Conv2D(16, 3, padding='same', activation='relu', name='bl_2')
base_layers = layers.MaxPooling2D(name='bl_3')(base_layers)
base_layers = layers.Conv2D(32, 3, padding='same', activation='relu', name='bl_4')
base_layers = layers.MaxPooling2D(name='bl_5')(base_layers)
base_layers = layers.Conv2D(64, 3, padding='same', activation='relu', name='bl_6')
base_layers = layers.MaxPooling2D(name='bl_7')(base_layers)
base_layers = layers.Flatten(name='bl_8')(base_layers)

locator_branch = layers.Dense(128, activation='relu', name='bb_1')(base_layers)
locator_branch = layers.Dense(64, activation='relu', name='bb_2')(locator_branch)
locator_branch = layers.Dense(32, activation='relu', name='bb_3')(locator_branch)
locator_branch = layers.Dense(4, activation='sigmoid', name='bb_head')(locator_branch)
```

In [6]:

```
model = tf.keras.Model(input_layer, outputs=[locator_branch])
```

In [7]:

```
model.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	[(None, 480, 480, 3)]	0
bl_1 (Rescaling)	(None, 480, 480, 3)	0
bl_2 (Conv2D)	(None, 480, 480, 16)	448
bl_3 (MaxPooling2D)	(None, 240, 240, 16)	0
bl_4 (Conv2D)	(None, 240, 240, 32)	4640
bl_5 (MaxPooling2D)	(None, 120, 120, 32)	0
bl_6 (Conv2D)	(None, 120, 120, 64)	18496
bl_7 (MaxPooling2D)	(None, 60, 60, 64)	0
bl_8 (Flatten)	(None, 230400)	0
bb_1 (Dense)	(None, 128)	29491328
bb_2 (Dense)	(None, 64)	8256
bb_3 (Dense)	(None, 32)	2080
bb_head (Dense)	(None, 4)	132
<hr/>		
Total params: 29525380 (112.63 MB)		
Trainable params: 29525380 (112.63 MB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [8]: losses = tf.keras.losses.MSE
```

```
In [9]: model.compile(loss=losses, optimizer='Adam', metrics=['accuracy'])
```

```
In [10]: training_epochs = 50
print(len(train_images))

train_images = np.array(train_images)
print(np.shape(train_images))
train_targets = np.array(train_targets)
train_labels = np.array(train_labels)

validation_images = np.array(validation_images)
validation_targets = np.array(validation_targets)
validation_labels = np.array(validation_labels)
```

```
100
(100, 480, 480, 3)
```

```
In [11]: trainTargets = {
```

```
        "bb_head": train_targets
    }

validationTargets = {
    "bb_head": validation_targets
}
```

```
In [12]: print(type(trainTargets))
```

```
<class 'dict'>
```

```
In [13]: history = model.fit(train_images, trainTargets,
                           validation_data=(validation_images, validationTargets),
                           batch_size=4,
                           epochs=training_epochs,
                           shuffle=True,
                           verbose=1)
```

Epoch 1/50
25/25 [=====] - 5s 167ms/step - loss: 0.0630 - accuracy: 0.5900 - val_loss: 0.0186 - val_accuracy: 0.5854
Epoch 2/50
25/25 [=====] - 4s 158ms/step - loss: 0.0154 - accuracy: 0.6400 - val_loss: 0.0169 - val_accuracy: 0.5854
Epoch 3/50
25/25 [=====] - 4s 156ms/step - loss: 0.0091 - accuracy: 0.8000 - val_loss: 0.0096 - val_accuracy: 0.7805
Epoch 4/50
25/25 [=====] - 4s 160ms/step - loss: 0.0040 - accuracy: 0.9200 - val_loss: 0.0092 - val_accuracy: 0.7805
Epoch 5/50
25/25 [=====] - 4s 158ms/step - loss: 0.0030 - accuracy: 0.8900 - val_loss: 0.0080 - val_accuracy: 0.8780
Epoch 6/50
25/25 [=====] - 4s 156ms/step - loss: 0.0014 - accuracy: 0.9100 - val_loss: 0.0074 - val_accuracy: 0.8049
Epoch 7/50
25/25 [=====] - 4s 169ms/step - loss: 7.8116e-04 - accuracy: 0.9400 - val_loss: 0.0084 - val_accuracy: 0.8049
Epoch 8/50
25/25 [=====] - 4s 166ms/step - loss: 7.3395e-04 - accuracy: 0.9500 - val_loss: 0.0077 - val_accuracy: 0.8293
Epoch 9/50
25/25 [=====] - 4s 160ms/step - loss: 6.4111e-04 - accuracy: 0.9500 - val_loss: 0.0074 - val_accuracy: 0.8293
Epoch 10/50
25/25 [=====] - 4s 173ms/step - loss: 5.2509e-04 - accuracy: 0.9500 - val_loss: 0.0076 - val_accuracy: 0.8049
Epoch 11/50
25/25 [=====] - 4s 162ms/step - loss: 3.8873e-04 - accuracy: 0.9800 - val_loss: 0.0077 - val_accuracy: 0.8293
Epoch 12/50
25/25 [=====] - 4s 164ms/step - loss: 3.8200e-04 - accuracy: 0.9500 - val_loss: 0.0074 - val_accuracy: 0.8537
Epoch 13/50
25/25 [=====] - 4s 157ms/step - loss: 3.3092e-04 - accuracy: 0.9900 - val_loss: 0.0076 - val_accuracy: 0.8293
Epoch 14/50
25/25 [=====] - 4s 155ms/step - loss: 3.2849e-04 - accuracy: 0.9700 - val_loss: 0.0075 - val_accuracy: 0.7561
Epoch 15/50
25/25 [=====] - 4s 156ms/step - loss: 3.2339e-04 - accuracy: 0.9400 - val_loss: 0.0072 - val_accuracy: 0.8049
Epoch 16/50
25/25 [=====] - 4s 155ms/step - loss: 2.6963e-04 - accuracy: 0.9600 - val_loss: 0.0074 - val_accuracy: 0.8293
Epoch 17/50
25/25 [=====] - 4s 161ms/step - loss: 3.0304e-04 - accuracy: 0.9700 - val_loss: 0.0075 - val_accuracy: 0.8293
Epoch 18/50
25/25 [=====] - 4s 157ms/step - loss: 3.4072e-04 - accuracy: 0.9500 - val_loss: 0.0076 - val_accuracy: 0.8293
Epoch 19/50
25/25 [=====] - 4s 155ms/step - loss: 4.4070e-04 - a

accuracy: 0.9700 - val_loss: 0.0080 - val_accuracy: 0.8293
Epoch 20/50
25/25 [=====] - 4s 155ms/step - loss: 4.9040e-04 - accuracy: 0.9600 - val_loss: 0.0074 - val_accuracy: 0.8293
Epoch 21/50
25/25 [=====] - 4s 157ms/step - loss: 2.3586e-04 - accuracy: 0.9700 - val_loss: 0.0078 - val_accuracy: 0.8293
Epoch 22/50
25/25 [=====] - 4s 157ms/step - loss: 2.3914e-04 - accuracy: 0.9900 - val_loss: 0.0076 - val_accuracy: 0.8293
Epoch 23/50
25/25 [=====] - 4s 158ms/step - loss: 2.3998e-04 - accuracy: 0.9800 - val_loss: 0.0073 - val_accuracy: 0.8537
Epoch 24/50
25/25 [=====] - 4s 159ms/step - loss: 1.6565e-04 - accuracy: 0.9700 - val_loss: 0.0074 - val_accuracy: 0.8293
Epoch 25/50
25/25 [=====] - 4s 164ms/step - loss: 1.5617e-04 - accuracy: 0.9700 - val_loss: 0.0072 - val_accuracy: 0.8537
Epoch 26/50
25/25 [=====] - 4s 158ms/step - loss: 1.6890e-04 - accuracy: 1.0000 - val_loss: 0.0073 - val_accuracy: 0.8293
Epoch 27/50
25/25 [=====] - 4s 157ms/step - loss: 1.2552e-04 - accuracy: 0.9800 - val_loss: 0.0072 - val_accuracy: 0.8537
Epoch 28/50
25/25 [=====] - 4s 164ms/step - loss: 1.0159e-04 - accuracy: 1.0000 - val_loss: 0.0074 - val_accuracy: 0.8293
Epoch 29/50
25/25 [=====] - 4s 164ms/step - loss: 1.0221e-04 - accuracy: 0.9900 - val_loss: 0.0075 - val_accuracy: 0.8537
Epoch 30/50
25/25 [=====] - 4s 160ms/step - loss: 1.3653e-04 - accuracy: 0.9800 - val_loss: 0.0075 - val_accuracy: 0.8293
Epoch 31/50
25/25 [=====] - 4s 161ms/step - loss: 1.3919e-04 - accuracy: 0.9800 - val_loss: 0.0075 - val_accuracy: 0.8537
Epoch 32/50
25/25 [=====] - 4s 156ms/step - loss: 1.1490e-04 - accuracy: 0.9700 - val_loss: 0.0074 - val_accuracy: 0.8293
Epoch 33/50
25/25 [=====] - 4s 160ms/step - loss: 1.2204e-04 - accuracy: 0.9900 - val_loss: 0.0072 - val_accuracy: 0.8537
Epoch 34/50
25/25 [=====] - 4s 157ms/step - loss: 1.2920e-04 - accuracy: 0.9900 - val_loss: 0.0073 - val_accuracy: 0.8293
Epoch 35/50
25/25 [=====] - 4s 156ms/step - loss: 1.4859e-04 - accuracy: 0.9800 - val_loss: 0.0079 - val_accuracy: 0.8049
Epoch 36/50
25/25 [=====] - 4s 159ms/step - loss: 1.4652e-04 - accuracy: 0.9600 - val_loss: 0.0073 - val_accuracy: 0.8293
Epoch 37/50
25/25 [=====] - 4s 156ms/step - loss: 1.9214e-04 - accuracy: 0.9900 - val_loss: 0.0073 - val_accuracy: 0.8537
Epoch 38/50

```
25/25 [=====] - 4s 158ms/step - loss: 1.5317e-04 - accuracy: 0.9700 - val_loss: 0.0074 - val_accuracy: 0.8537
Epoch 39/50
25/25 [=====] - 4s 155ms/step - loss: 1.5015e-04 - accuracy: 0.9700 - val_loss: 0.0074 - val_accuracy: 0.8537
Epoch 40/50
25/25 [=====] - 4s 156ms/step - loss: 1.3490e-04 - accuracy: 0.9900 - val_loss: 0.0076 - val_accuracy: 0.8537
Epoch 41/50
25/25 [=====] - 4s 161ms/step - loss: 1.4752e-04 - accuracy: 0.9900 - val_loss: 0.0076 - val_accuracy: 0.8293
Epoch 42/50
25/25 [=====] - 4s 166ms/step - loss: 1.8806e-04 - accuracy: 0.9900 - val_loss: 0.0077 - val_accuracy: 0.7073
Epoch 43/50
25/25 [=====] - 4s 161ms/step - loss: 3.2610e-04 - accuracy: 0.9600 - val_loss: 0.0076 - val_accuracy: 0.8293
Epoch 44/50
25/25 [=====] - 4s 155ms/step - loss: 4.4658e-04 - accuracy: 0.9700 - val_loss: 0.0074 - val_accuracy: 0.8537
Epoch 45/50
25/25 [=====] - 4s 154ms/step - loss: 3.1221e-04 - accuracy: 0.9800 - val_loss: 0.0079 - val_accuracy: 0.8780
Epoch 46/50
25/25 [=====] - 4s 156ms/step - loss: 3.0187e-04 - accuracy: 0.9600 - val_loss: 0.0081 - val_accuracy: 0.8293
Epoch 47/50
25/25 [=====] - 4s 153ms/step - loss: 3.8423e-04 - accuracy: 0.9900 - val_loss: 0.0076 - val_accuracy: 0.8293
Epoch 48/50
25/25 [=====] - 4s 151ms/step - loss: 3.9240e-04 - accuracy: 0.9400 - val_loss: 0.0077 - val_accuracy: 0.8293
Epoch 49/50
25/25 [=====] - 4s 153ms/step - loss: 4.2223e-04 - accuracy: 0.9700 - val_loss: 0.0076 - val_accuracy: 0.8293
Epoch 50/50
25/25 [=====] - 4s 156ms/step - loss: 2.7108e-04 - accuracy: 0.9800 - val_loss: 0.0074 - val_accuracy: 0.8049
```

```
In [14]: bb_accuracy = history.history['accuracy']
bb_val_acc = history.history['val_accuracy']

bb_loss = history.history['loss']
bb_val_loss = history.history['val_loss']

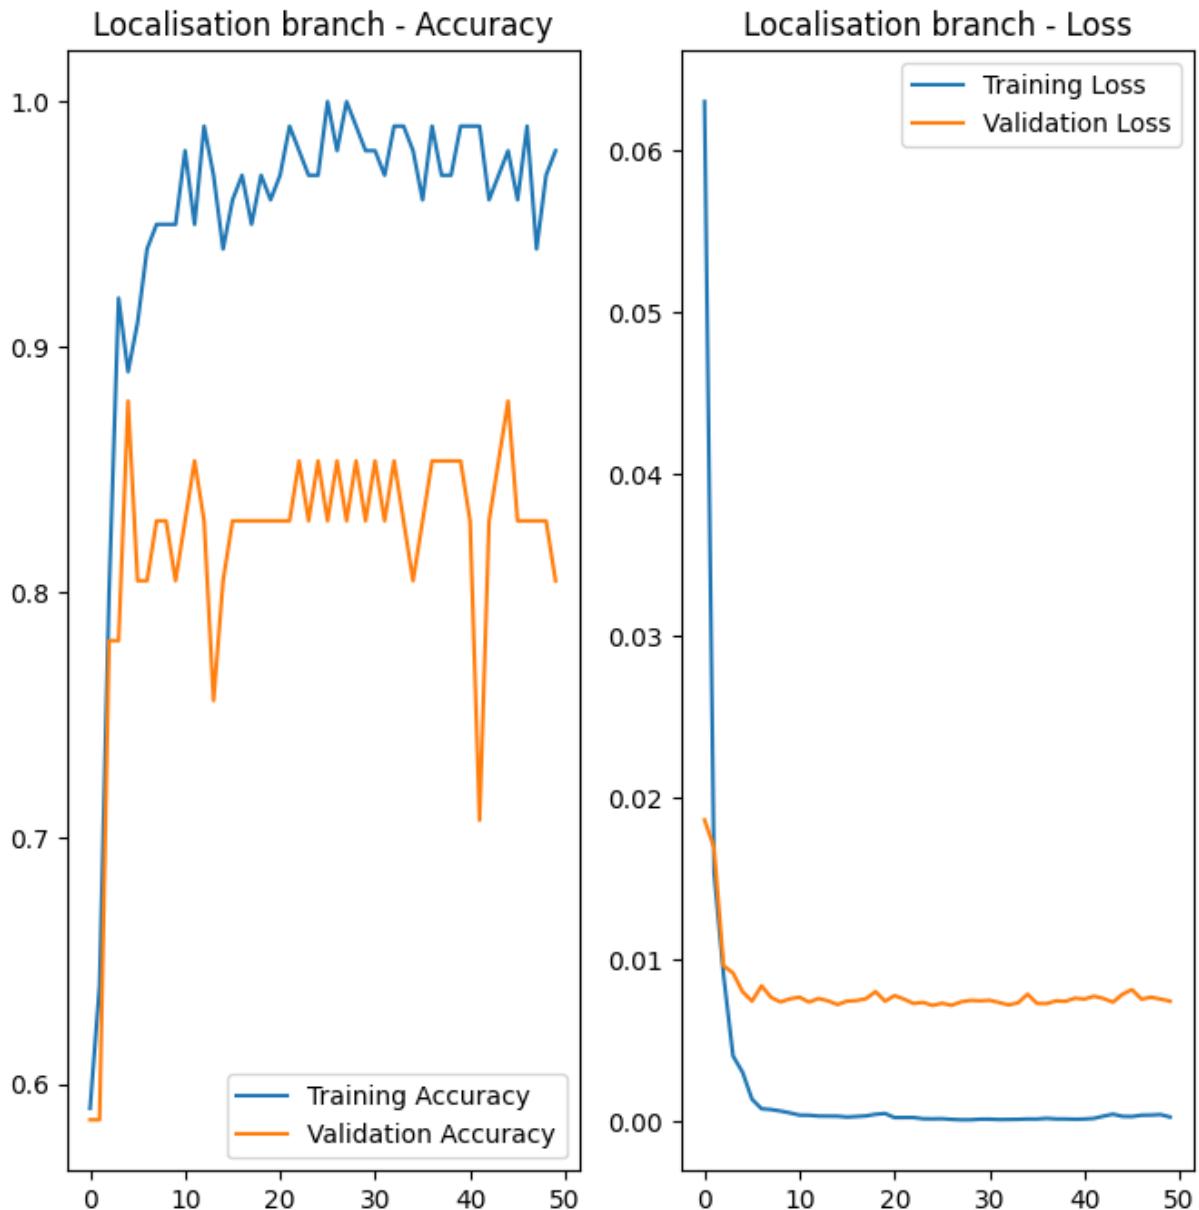
epochs_range = range(training_epochs)

plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(epochs_range, bb_accuracy, label='Training Accuracy')
plt.plot(epochs_range, bb_val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Localisation branch - Accuracy')
```

```

plt.subplot(1, 2, 2)
plt.plot(epochs_range, bb_loss, label='Training Loss')
plt.plot(epochs_range, bb_val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Localisation branch - Loss')
plt.show()

```



```

In [19]: test_dir = pathlib.Path('Images/test')
img_paths = list(test_dir.glob('*.*'))

plt.figure(figsize=(20, 20))

for i, img_path in enumerate(img_paths):
    test_img = keras.preprocessing.image.load_img(img_path, target_size=(hei
    img_arr = keras.preprocessing.image.img_to_array(test_img)
    img_arr = tf.expand_dims(img_arr, 0)

    predictions = model.predict(img_arr)

```

```
print(predictions)

bbox = predictions[0]
bbox = [bbox[0] * width, bbox[1] * height, bbox[2] * width, bbox[3] * he

draw1 = Draw(test_img)
draw1.rectangle(bbox, outline='cyan', width=2)
#test_img

ax = plt.subplot(6, 5, i+1)
plt.imshow(test_img)
plt.axis("off")
```

1/1 [=====] - 0s 25ms/step
[[0.46281928 0.4482857 0.64412576 0.5189963]]
1/1 [=====] - 0s 22ms/step
[[0.45809156 0.43809587 0.64807606 0.5175914]]
1/1 [=====] - 0s 22ms/step
[[0.4651276 0.4345678 0.5910957 0.50153315]]
1/1 [=====] - 0s 22ms/step
[[0.4927985 0.4561605 0.6270964 0.5141671]]
1/1 [=====] - 0s 20ms/step
[[0.3678292 0.1978642 0.67193484 0.40086251]]
1/1 [=====] - 0s 21ms/step
[[0.5134349 0.48509794 0.6603877 0.52138096]]
1/1 [=====] - 0s 20ms/step
[[0.52950895 0.45523557 0.6484508 0.50636464]]
1/1 [=====] - 0s 21ms/step
[[0.24008419 0.17058705 0.62376064 0.42314884]]
1/1 [=====] - 0s 21ms/step
[[0.2577669 0.27012616 0.668748 0.45836544]]
1/1 [=====] - 0s 21ms/step
[[0.41723317 0.42936638 0.7196637 0.55487907]]
1/1 [=====] - 0s 21ms/step
[[0.24884035 0.3457871 0.5452062 0.49743208]]
1/1 [=====] - 0s 21ms/step
[[0.40952894 0.38085672 0.6494406 0.49572852]]
1/1 [=====] - 0s 21ms/step
[[0.5340578 0.4925578 0.6788949 0.5500652]]
1/1 [=====] - 0s 22ms/step
[[0.3045203 0.1668789 0.5639035 0.37595075]]
1/1 [=====] - 0s 23ms/step
[[0.3875958 0.4362633 0.6177927 0.5619406]]
1/1 [=====] - 0s 22ms/step
[[0.38838744 0.41352844 0.5455041 0.5310311]]
1/1 [=====] - 0s 22ms/step
[[0.3652346 0.30933797 0.6295046 0.48354056]]
1/1 [=====] - 0s 20ms/step
[[0.34637856 0.37179923 0.7962475 0.5659767]]
1/1 [=====] - 0s 21ms/step
[[0.48236397 0.37327507 0.65130836 0.46984568]]
1/1 [=====] - 0s 22ms/step
[[0.22652459 0.46971732 0.72277296 0.6946776]]
1/1 [=====] - 0s 22ms/step
[[0.29687464 0.5012101 0.6083357 0.65008557]]
1/1 [=====] - 0s 21ms/step
[[0.3219837 0.2806773 0.58761704 0.4613152]]
1/1 [=====] - 0s 21ms/step
[[0.2716597 0.39990032 0.5207666 0.5366165]]
1/1 [=====] - 0s 21ms/step
[[0.37162805 0.3941619 0.5627825 0.50594443]]
1/1 [=====] - 0s 23ms/step
[[0.2327291 0.49224025 0.6758387 0.7111752]]
1/1 [=====] - 0s 21ms/step
[[0.27799928 0.32281303 0.53691804 0.5130597]]
1/1 [=====] - 0s 21ms/step
[[0.3106796 0.50874597 0.6331433 0.6627335]]
1/1 [=====] - 0s 22ms/step
[[0.5508148 0.46227247 0.71086633 0.5262203]]

```
1/1 [=====] - 0s 21ms/step
[[0.43896824 0.3708446 0.5743124 0.4768832 ]]
1/1 [=====] - 0s 21ms/step
[[0.3138904 0.3374242 0.59048593 0.5068008 ]]
```



```
In [18]: test_dir = pathlib.Path('Images/test2')
img_paths = list(test_dir.glob('*.*jpg'))

plt.figure(figsize=(20, 20))

for i, img_path in enumerate(img_paths):
    test_img = keras.preprocessing.image.load_img(img_path, target_size=(height, width))
    img_arr = keras.preprocessing.image.img_to_array(test_img)
    img_arr = tf.expand_dims(img_arr, 0)

    predictions = model.predict(img_arr)
    print(predictions)
```

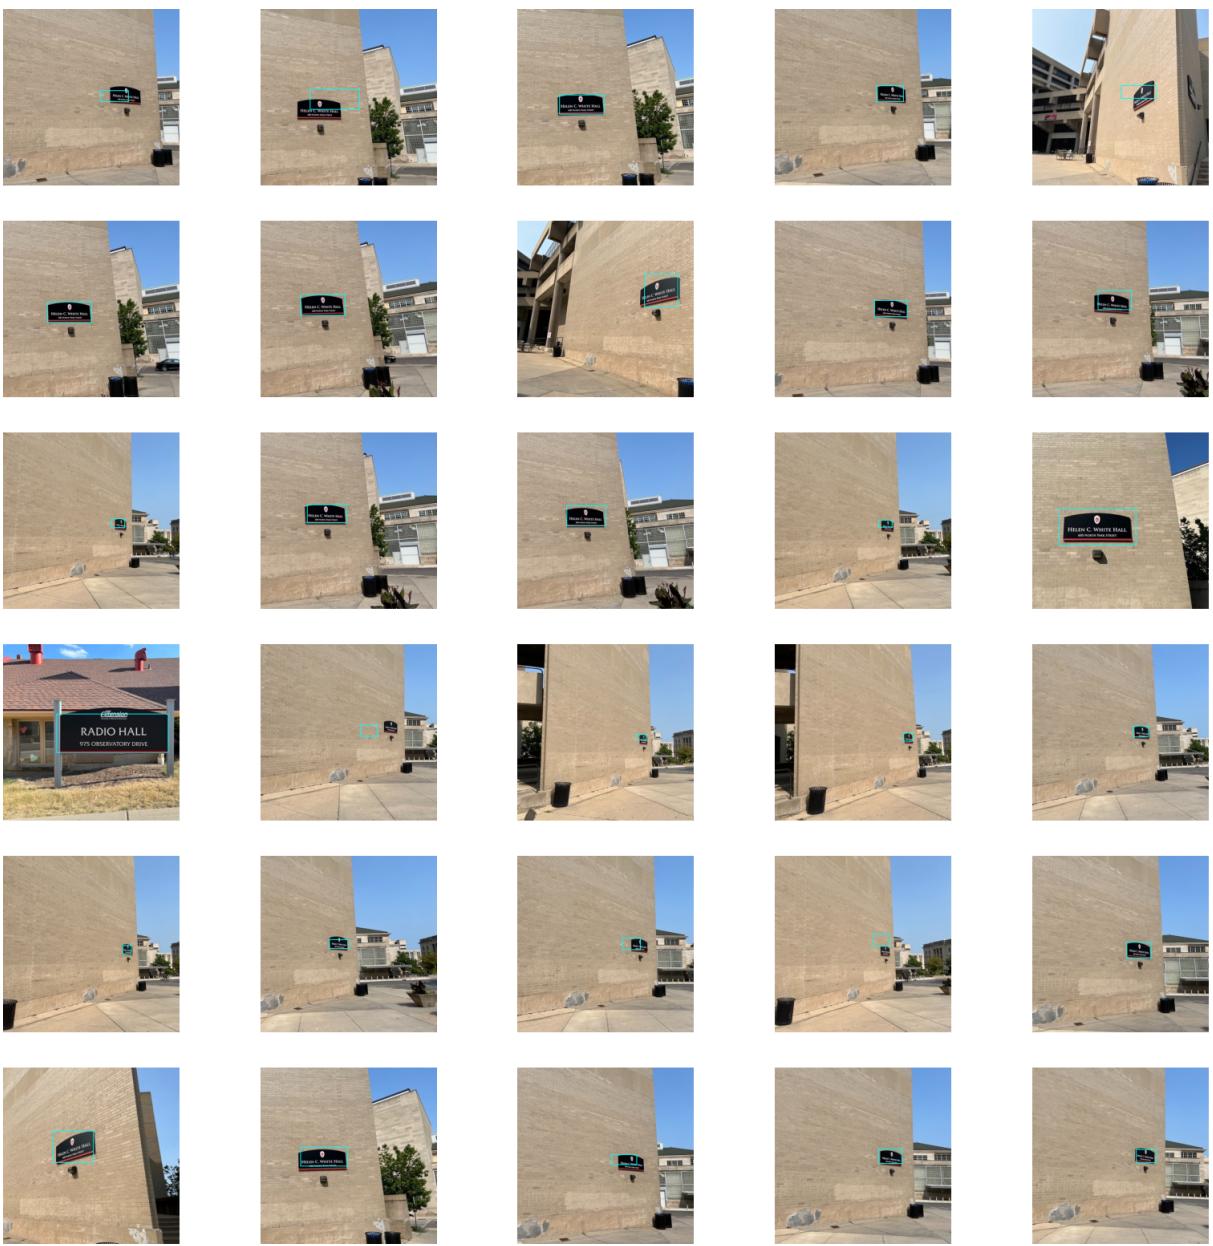
```
bbox = predictions[0]
bbox = [bbox[0] * width, bbox[1] * height, bbox[2] * width, bbox[3] * he

draw1 = Draw(test_img)
draw1.rectangle(bbox, outline='cyan', width=2)
#test_img

ax = plt.subplot(6, 5, i+1)
plt.imshow(test_img)
plt.axis("off")
```

1/1 [=====] - 0s 24ms/step
[[0.5508148 0.46227247 0.71086633 0.5262203]]
1/1 [=====] - 0s 21ms/step
[[0.27948442 0.45392978 0.5624292 0.5707499]]
1/1 [=====] - 0s 23ms/step
[[0.23795603 0.48662624 0.49212334 0.597246]]
1/1 [=====] - 0s 21ms/step
[[0.5766924 0.4337504 0.73119575 0.5269951]]
1/1 [=====] - 0s 21ms/step
[[0.5021373 0.432592 0.7046188 0.507299]]
1/1 [=====] - 0s 22ms/step
[[0.25016764 0.45876962 0.5029787 0.5773083]]
1/1 [=====] - 0s 22ms/step
[[0.22644201 0.41780722 0.47954834 0.538838]]
1/1 [=====] - 0s 22ms/step
[[0.7196603 0.300886 0.91185826 0.48843247]]
1/1 [=====] - 0s 20ms/step
[[0.56389827 0.45096198 0.75064075 0.5453222]]
1/1 [=====] - 0s 21ms/step
[[0.37162805 0.3941619 0.5627825 0.50594443]]
1/1 [=====] - 0s 21ms/step
[[0.61357653 0.48915058 0.6906426 0.53975666]]
1/1 [=====] - 0s 22ms/step
[[0.26372528 0.4091444 0.48885912 0.51359975]]
1/1 [=====] - 0s 22ms/step
[[0.2808542 0.4140245 0.50214106 0.5298429]]
1/1 [=====] - 0s 23ms/step
[[0.5884565 0.49607006 0.66516316 0.5386007]]
1/1 [=====] - 0s 22ms/step
[[0.15064299 0.43101126 0.589345 0.6377273]]
1/1 [=====] - 0s 30ms/step
[[0.30288565 0.3947372 0.9362964 0.61624897]]
1/1 [=====] - 0s 25ms/step
[[0.56475455 0.45758483 0.66402054 0.52319425]]
1/1 [=====] - 0s 22ms/step
[[0.6728495 0.5120985 0.7318277 0.5422966]]
1/1 [=====] - 0s 28ms/step
[[0.72825605 0.501818 0.78219754 0.55108356]]
1/1 [=====] - 0s 21ms/step
[[0.569398 0.46273518 0.65873605 0.52646035]]
1/1 [=====] - 0s 24ms/step
[[0.67778647 0.50654614 0.7285546 0.5578932]]
1/1 [=====] - 0s 22ms/step
[[0.38943353 0.47383836 0.5040472 0.5269981]]
1/1 [=====] - 0s 21ms/step
[[0.59510404 0.46457112 0.6982028 0.53145826]]
1/1 [=====] - 0s 21ms/step
[[0.55704635 0.44272104 0.6525481 0.5143303]]
1/1 [=====] - 0s 20ms/step
[[0.5339154 0.4883845 0.67312926 0.58119315]]
1/1 [=====] - 0s 20ms/step
[[0.28143165 0.35995865 0.51439726 0.538551]]
1/1 [=====] - 0s 22ms/step
[[0.22698306 0.44876617 0.49433452 0.55972886]]
1/1 [=====] - 0s 21ms/step
[[0.5340578 0.4925578 0.6788949 0.5500652]]

```
1/1 [=====] - 0s 23ms/step  
[[0.5870225 0.46006936 0.7226593 0.5456031 ]]  
1/1 [=====] - 0s 22ms/step  
[[0.58754927 0.4627597 0.70279837 0.5408365 ]]
```



```
In [20]: test_dir = pathlib.Path('Images/test3')  
img_paths = list(test_dir.glob('*.*'))  
  
plt.figure(figsize=(20, 20))  
  
for i, img_path in enumerate(img_paths):  
    test_img = keras.preprocessing.image.load_img(img_path, target_size=(hei  
img_arr = keras.preprocessing.image.img_to_array(test_img)  
img_arr = tf.expand_dims(img_arr, 0)  
  
predictions = model.predict(img_arr)  
print(predictions)
```

```
bbox = predictions[0]
bbox = [bbox[0] * width, bbox[1] * height, bbox[2] * width, bbox[3] * height]

draw1 = Draw(test_img)
draw1.rectangle(bbox, outline='cyan', width=2)
#test_img

ax = plt.subplot(6, 5, i+1)
plt.imshow(test_img)
plt.axis("off")
```

1/1 [=====] - 0s 25ms/step
[[0.27562454 0.51960427 0.76914907 0.7453672]]
1/1 [=====] - 0s 21ms/step
[[0.3427732 0.40936208 0.6974351 0.6188924]]
1/1 [=====] - 0s 21ms/step
[[0.16135658 0.4131523 0.58814675 0.62416255]]
1/1 [=====] - 0s 23ms/step
[[0.05002226 0.30050468 0.5418818 0.63023776]]
1/1 [=====] - 0s 20ms/step
[[0.2646662 0.46822667 0.64990705 0.64745367]]
1/1 [=====] - 0s 20ms/step
[[0.48236397 0.37327507 0.65130836 0.46984568]]
1/1 [=====] - 0s 21ms/step
[[0.4635744 0.40470418 0.6127713 0.48057014]]
1/1 [=====] - 0s 21ms/step
[[0.3106796 0.50874597 0.6331433 0.6627335]]
1/1 [=====] - 0s 21ms/step
[[0.29687464 0.5012101 0.6083357 0.65008557]]
1/1 [=====] - 0s 21ms/step
[[0.5011715 0.36606988 0.606852 0.43000513]]
1/1 [=====] - 0s 21ms/step
[[0.47196296 0.40942663 0.5799865 0.49884665]]
1/1 [=====] - 0s 21ms/step
[[0.42158067 0.42885673 0.51993716 0.49721068]]
1/1 [=====] - 0s 22ms/step
[[0.35445973 0.45656577 0.5844166 0.5606887]]
1/1 [=====] - 0s 23ms/step
[[0.0377108 0.2653488 0.92587966 0.6370258]]
1/1 [=====] - 0s 23ms/step
[[0.27335805 0.5359581 0.59752095 0.74745506]]
1/1 [=====] - 0s 22ms/step
[[0.32500944 0.21977612 0.6413764 0.45205694]]
1/1 [=====] - 0s 22ms/step
[[0.14872396 0.39000696 0.5304915 0.6055774]]
1/1 [=====] - 0s 21ms/step
[[0.300133 0.4062915 0.67116624 0.601342]]
1/1 [=====] - 0s 21ms/step
[[0.27226096 0.40797213 0.52191436 0.5422552]]
1/1 [=====] - 0s 25ms/step
[[0.38838744 0.41352844 0.5455041 0.5310311]]
1/1 [=====] - 0s 22ms/step
[[0.19434287 0.3881554 0.5858703 0.6046384]]
1/1 [=====] - 0s 21ms/step
[[0.44096434 0.43422946 0.7089658 0.5315649]]
1/1 [=====] - 0s 21ms/step
[[0.15326469 0.14976402 0.8252921 0.43398252]]
1/1 [=====] - 0s 21ms/step
[[0.362608 0.3997101 0.42381474 0.48810494]]
1/1 [=====] - 0s 20ms/step
[[0.3534399 0.41195378 0.53806984 0.5258205]]
1/1 [=====] - 0s 21ms/step
[[0.32357442 0.5424925 0.5683681 0.76464134]]
1/1 [=====] - 0s 23ms/step
[[0.2716597 0.39990032 0.5207666 0.5366165]]
1/1 [=====] - 0s 22ms/step
[[0.35690656 0.39211142 0.8712833 0.7201662]]

```
1/1 [=====] - 0s 23ms/step
[[0.40456724 0.5567611 0.79680705 0.7988434 ]]
1/1 [=====] - 0s 22ms/step
[[0.27190286 0.38744965 0.51739466 0.531398 ]]
```



```
In [21]: test_dir = pathlib.Path('Images/test4')
img_paths = list(test_dir.glob('*.*jpg'))

plt.figure(figsize=(20, 20))

for i, img_path in enumerate(img_paths):
    test_img = keras.preprocessing.image.load_img(img_path, target_size=(height, width))
    img_arr = keras.preprocessing.image.img_to_array(test_img)
    img_arr = tf.expand_dims(img_arr, 0)

    predictions = model.predict(img_arr)
    print(predictions)
```

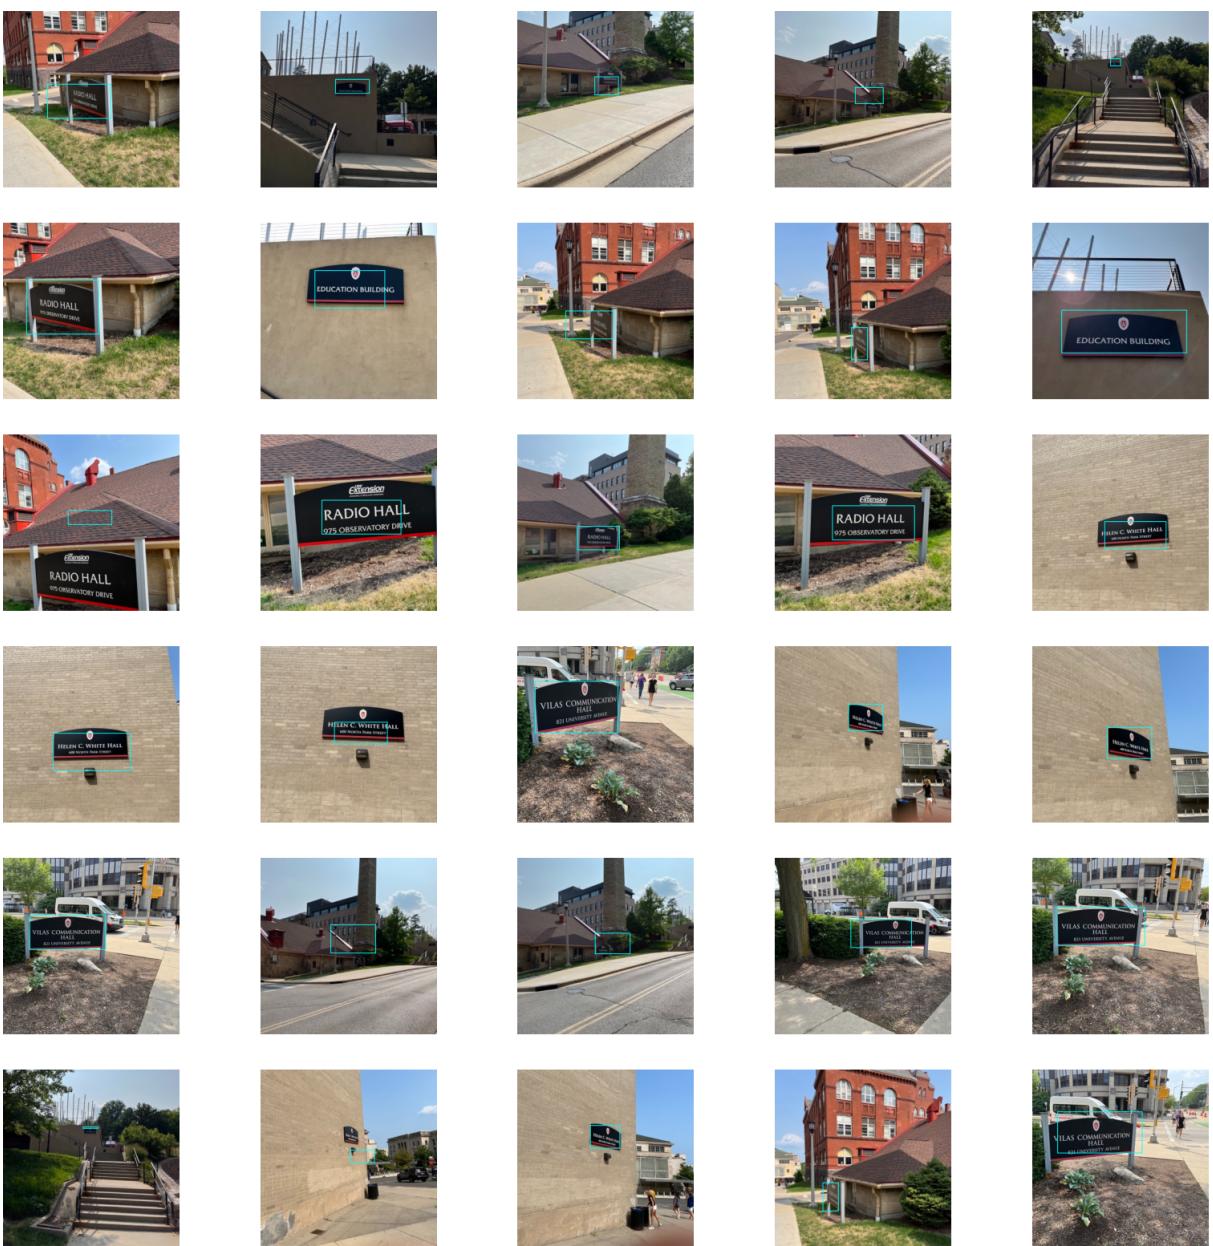
```
bbox = predictions[0]
bbox = [bbox[0] * width, bbox[1] * height, bbox[2] * width, bbox[3] * height]

draw1 = Draw(test_img)
draw1.rectangle(bbox, outline='cyan', width=2)
#test_img

ax = plt.subplot(6, 5, i+1)
plt.imshow(test_img)
plt.axis("off")
```

1/1 [=====] - 0s 24ms/step
[[0.2510562 0.4137792 0.6104995 0.6087847]]
1/1 [=====] - 0s 23ms/step
[[0.4235955 0.38619325 0.6166845 0.4671047]]
1/1 [=====] - 0s 21ms/step
[[0.43896824 0.3708446 0.5743124 0.4768832]]
1/1 [=====] - 0s 21ms/step
[[0.4546889 0.4315351 0.61661077 0.52640486]]
1/1 [=====] - 0s 21ms/step
[[0.44193968 0.26979816 0.50061476 0.3123401]]
1/1 [=====] - 0s 22ms/step
[[0.13265258 0.32318234 0.54286706 0.6323796]]
1/1 [=====] - 0s 21ms/step
[[0.3074138 0.27073023 0.70546246 0.48378888]]
1/1 [=====] - 0s 22ms/step
[[0.27490455 0.49958822 0.5489544 0.6595868]]
1/1 [=====] - 0s 21ms/step
[[0.43471214 0.5918256 0.52591103 0.7780924]]
1/1 [=====] - 0s 21ms/step
[[0.16464922 0.49580422 0.87658507 0.7375429]]
1/1 [=====] - 0s 21ms/step
[[0.36771014 0.42996195 0.61807287 0.5121916]]
1/1 [=====] - 0s 20ms/step
[[0.34637856 0.37179923 0.7962475 0.5659767]]
1/1 [=====] - 0s 21ms/step
[[0.34122595 0.52074367 0.5807272 0.65524465]]
1/1 [=====] - 0s 20ms/step
[[0.32584146 0.40360156 0.7925955 0.5999362]]
1/1 [=====] - 0s 21ms/step
[[0.4102288 0.4936819 0.75883186 0.64611745]]
1/1 [=====] - 0s 22ms/step
[[0.28509676 0.49246094 0.72467595 0.7076998]]
1/1 [=====] - 0s 23ms/step
[[0.41723317 0.42936638 0.7196637 0.55487907]]
1/1 [=====] - 0s 21ms/step
[[0.10070655 0.1993469 0.572769 0.49313492]]
1/1 [=====] - 0s 20ms/step
[[0.4169833 0.32767287 0.6140481 0.48056242]]
1/1 [=====] - 0s 22ms/step
[[0.42469275 0.45854744 0.67899555 0.6386448]]
1/1 [=====] - 0s 22ms/step
[[0.15188938 0.32100046 0.58275294 0.5196867]]
1/1 [=====] - 0s 21ms/step
[[0.39707357 0.37734184 0.65339696 0.54004484]]
1/1 [=====] - 0s 21ms/step
[[0.44108802 0.42685682 0.6404031 0.54515934]]
1/1 [=====] - 0s 21ms/step
[[0.43056828 0.3485165 0.77590543 0.50858355]]
1/1 [=====] - 0s 20ms/step
[[0.13379371 0.28053227 0.6385824 0.49647635]]
1/1 [=====] - 0s 20ms/step
[[0.45289746 0.32408655 0.540896 0.33418804]]
1/1 [=====] - 0s 21ms/step
[[0.5065976 0.45509338 0.646822 0.52911615]]
1/1 [=====] - 0s 22ms/step
[[0.41429737 0.31618136 0.58271533 0.44585806]]

```
1/1 [=====] - 0s 22ms/step
[[0.27150837 0.6410867 0.3619813 0.8036262 ]]
1/1 [=====] - 0s 21ms/step
[[0.14276077 0.23916848 0.6249135 0.47387236]]
```



```
In [22]: test_dir = pathlib.Path('Images/test5')
img_paths = list(test_dir.glob('*.*jpg'))

plt.figure(figsize=(20, 20))

for i, img_path in enumerate(img_paths):
    test_img = keras.preprocessing.image.load_img(img_path, target_size=(height, width))
    img_arr = keras.preprocessing.image.img_to_array(test_img)
    img_arr = tf.expand_dims(img_arr, 0)

    predictions = model.predict(img_arr)
    print(predictions)
```

```
bbox = predictions[0]
bbox = [bbox[0] * width, bbox[1] * height, bbox[2] * width, bbox[3] * height]

draw1 = Draw(test_img)
draw1.rectangle(bbox, outline='cyan', width=2)
#test_img

ax = plt.subplot(6, 5, i+1)
plt.imshow(test_img)
plt.axis("off")
```

1/1 [=====] - 0s 22ms/step
[[0.33223802 0.59642285 0.5588345 0.70999163]]
1/1 [=====] - 0s 22ms/step
[[0.3675995 0.58751374 0.57128257 0.70018375]]
1/1 [=====] - 0s 21ms/step
[[0.4237034 0.46278632 0.6399201 0.5874194]]
1/1 [=====] - 0s 21ms/step
[[0.23885894 0.5395471 0.68205005 0.7528144]]
1/1 [=====] - 0s 20ms/step
[[0.36012992 0.45259076 0.5929761 0.5779837]]
1/1 [=====] - 0s 21ms/step
[[0.3875958 0.4362633 0.6177927 0.5619406]]
1/1 [=====] - 0s 21ms/step
[[0.37523946 0.6474228 0.53258884 0.72801536]]
1/1 [=====] - 0s 24ms/step
[[0.3045203 0.1668789 0.5639035 0.37595075]]
1/1 [=====] - 0s 20ms/step
[[0.3678292 0.1978642 0.67193484 0.40086251]]
1/1 [=====] - 0s 20ms/step
[[0.29797414 0.5170383 0.54529923 0.64228535]]
1/1 [=====] - 0s 22ms/step
[[0.28343782 0.51995844 0.51164806 0.6388029]]
1/1 [=====] - 0s 20ms/step
[[0.20106408 0.30971932 0.5635595 0.47781274]]
1/1 [=====] - 0s 20ms/step
[[0.2704828 0.1510944 0.611092 0.39982355]]
1/1 [=====] - 0s 21ms/step
[[0.2918726 0.5422186 0.56276655 0.6653347]]
1/1 [=====] - 0s 23ms/step
[[0.37064472 0.18047646 0.7691673 0.35927618]]
1/1 [=====] - 0s 23ms/step
[[0.29021102 0.15173197 0.740669 0.36299196]]
1/1 [=====] - 0s 22ms/step
[[0.17506792 0.4739464 0.48851663 0.6256444]]
1/1 [=====] - 0s 21ms/step
[[0.26169574 0.20288815 0.56908613 0.4264687]]
1/1 [=====] - 0s 21ms/step
[[0.18574457 0.05897149 0.66875285 0.39790234]]
1/1 [=====] - 0s 22ms/step
[[0.17245556 0.43035227 0.6041787 0.6466182]]
1/1 [=====] - 0s 21ms/step
[[0.24008419 0.17058705 0.62376064 0.42314884]]
1/1 [=====] - 0s 22ms/step
[[0.28219068 0.12538248 0.63712496 0.37756842]]
1/1 [=====] - 0s 22ms/step
[[0.24884035 0.3457871 0.5452062 0.49743208]]
1/1 [=====] - 0s 21ms/step
[[0.18829297 0.45852417 0.55081743 0.62568253]]
1/1 [=====] - 0s 23ms/step
[[0.191195 0.10952079 0.67063695 0.42874798]]
1/1 [=====] - 0s 21ms/step
[[0.24957253 0.16560918 0.7232698 0.42074484]]
1/1 [=====] - 0s 21ms/step
[[0.37475926 0.368758 0.6742454 0.51245457]]
1/1 [=====] - 0s 22ms/step
[[0.35140514 0.15997148 0.7068531 0.45858818]]

```
1/1 [=====] - 0s 23ms/step  
[[0.32535937 0.26131204 0.7143306 0.45321685]]  
1/1 [=====] - 0s 22ms/step  
[[0.29776475 0.30409893 0.7900158 0.48628432]]
```



```
In [23]: test_dir = pathlib.Path('Images/test6')  
img_paths = list(test_dir.glob('*.*'))  
  
plt.figure(figsize=(20, 20))  
  
for i, img_path in enumerate(img_paths):  
    test_img = keras.preprocessing.image.load_img(img_path, target_size=(hei  
    img_arr = keras.preprocessing.image.img_to_array(test_img)  
    img_arr = tf.expand_dims(img_arr, 0)  
  
    predictions = model.predict(img_arr)  
    print(predictions)
```

```
bbox = predictions[0]
bbox = [bbox[0] * width, bbox[1] * height, bbox[2] * width, bbox[3] * height]

draw1 = Draw(test_img)
draw1.rectangle(bbox, outline='cyan', width=2)
#test_img

ax = plt.subplot(6, 5, i+1)
plt.imshow(test_img)
plt.axis("off")
```

1/1 [=====] - 0s 22ms/step
[[0.2327291 0.49224025 0.6758387 0.7111752]]
1/1 [=====] - 0s 22ms/step
[[0.20829114 0.4550381 0.67430377 0.679453]]
1/1 [=====] - 0s 20ms/step
[[0.3116277 0.5159067 0.60909283 0.6567496]]
1/1 [=====] - 0s 22ms/step
[[0.40655857 0.53224236 0.6991336 0.6928544]]
1/1 [=====] - 0s 20ms/step
[[0.22652459 0.46971732 0.72277296 0.6946776]]
1/1 [=====] - 0s 21ms/step
[[0.447338 0.41562882 0.6397395 0.5252299]]
1/1 [=====] - 0s 21ms/step
[[0.38051912 0.37776068 0.6190316 0.50985086]]
1/1 [=====] - 0s 22ms/step
[[0.40348735 0.2992174 0.6130713 0.5062635]]
1/1 [=====] - 0s 21ms/step
[[0.5063113 0.47302824 0.6295507 0.5510796]]
1/1 [=====] - 0s 21ms/step
[[0.5431022 0.1553719 0.6701129 0.24248819]]
1/1 [=====] - 0s 21ms/step
[[0.36141995 0.43957457 0.6358401 0.56771]]
1/1 [=====] - 0s 21ms/step
[[0.37903833 0.48783568 0.66656053 0.61499786]]
1/1 [=====] - 0s 21ms/step
[[0.31487933 0.31493482 0.5479151 0.5136556]]
1/1 [=====] - 0s 21ms/step
[[0.31223208 0.3220937 0.5832852 0.5386773]]
1/1 [=====] - 0s 21ms/step
[[0.40226215 0.47321102 0.62769467 0.57669985]]
1/1 [=====] - 0s 21ms/step
[[0.25933856 0.49691162 0.6510583 0.6902875]]
1/1 [=====] - 0s 22ms/step
[[0.3137554 0.55411386 0.5729551 0.68811035]]
1/1 [=====] - 0s 21ms/step
[[0.42388058 0.4254233 0.60392374 0.5423428]]
1/1 [=====] - 0s 21ms/step
[[0.45372295 0.41087252 0.61175925 0.5240063]]
1/1 [=====] - 0s 21ms/step
[[0.26279107 0.56249434 0.491068 0.7000996]]
1/1 [=====] - 0s 20ms/step
[[0.2715998 0.5455548 0.5926036 0.7566904]]
1/1 [=====] - 0s 24ms/step
[[0.27725953 0.4809487 0.6444552 0.65112174]]
1/1 [=====] - 0s 21ms/step
[[0.4303455 0.43262598 0.66636163 0.5441271]]
1/1 [=====] - 0s 21ms/step
[[0.3775917 0.48036537 0.6998768 0.6233284]]
1/1 [=====] - 0s 21ms/step
[[0.43839988 0.26942062 0.7315676 0.39840463]]
1/1 [=====] - 0s 21ms/step
[[0.43161765 0.32296252 0.6498027 0.43224645]]
1/1 [=====] - 0s 21ms/step
[[0.3652346 0.30933797 0.6295046 0.48354056]]
1/1 [=====] - 0s 21ms/step
[[0.32782176 0.5228783 0.6363258 0.68263274]]

```
1/1 [=====] - 0s 21ms/step  
[[0.3855456  0.37322435 0.6015159  0.5206019 ]]  
1/1 [=====] - 0s 30ms/step  
[[0.49604124 0.39397928 0.6233331  0.48209623]]
```



```
In [24]: test_dir = pathlib.Path('Images/test7')  
img_paths = list(test_dir.glob('*.*'))  
  
plt.figure(figsize=(20, 20))  
  
for i, img_path in enumerate(img_paths):  
    test_img = keras.preprocessing.image.load_img(img_path, target_size=(hei  
img_arr = keras.preprocessing.image.img_to_array(test_img)  
img_arr = tf.expand_dims(img_arr, 0)  
  
predictions = model.predict(img_arr)  
print(predictions)
```

```
bbox = predictions[0]
bbox = [bbox[0] * width, bbox[1] * height, bbox[2] * width, bbox[3] * height]

draw1 = Draw(test_img)
draw1.rectangle(bbox, outline='cyan', width=2)
#test_img

ax = plt.subplot(6, 5, i+1)
plt.imshow(test_img)
plt.axis("off")
```

1/1 [=====] - 0s 26ms/step
[[0.2761278 0.5055767 0.48461014 0.61189854]]
1/1 [=====] - 0s 26ms/step
[[0.13974635 0.21307302 0.75858974 0.49766526]]
1/1 [=====] - 0s 20ms/step
[[0.1966689 0.32284376 0.74078894 0.5515433]]
1/1 [=====] - 0s 21ms/step
[[0.36242515 0.42062062 0.6524133 0.5611263]]
1/1 [=====] - 0s 20ms/step
[[0.22706397 0.32426634 0.61815417 0.50860494]]
1/1 [=====] - 0s 21ms/step
[[0.29666585 0.37564605 0.72610617 0.5736548]]
1/1 [=====] - 0s 20ms/step
[[0.34970286 0.32000878 0.57606405 0.46727005]]
1/1 [=====] - 0s 22ms/step
[[0.3138904 0.3374242 0.59048593 0.5068008]]
1/1 [=====] - 0s 22ms/step
[[0.3086381 0.2571434 0.6690693 0.48323816]]
1/1 [=====] - 0s 20ms/step
[[0.700441 0.45190513 0.8020227 0.5788556]]
1/1 [=====] - 0s 23ms/step
[[0.7763053 0.4280716 0.9097229 0.5252449]]
1/1 [=====] - 0s 21ms/step
[[0.28498545 0.3396767 0.61296535 0.5038487]]
1/1 [=====] - 0s 23ms/step
[[0.2577669 0.27012616 0.668748 0.45836544]]
1/1 [=====] - 0s 30ms/step
[[0.3529035 0.2593301 0.8087755 0.42890564]]
1/1 [=====] - 0s 21ms/step
[[0.53780156 0.4972745 0.81704575 0.6151823]]
1/1 [=====] - 0s 20ms/step
[[0.30372196 0.26369703 0.6610698 0.45562196]]
1/1 [=====] - 0s 20ms/step
[[0.40952894 0.38085672 0.6494406 0.49572852]]
1/1 [=====] - 0s 20ms/step
[[0.6829286 0.3240939 0.86163557 0.44125754]]
1/1 [=====] - 0s 22ms/step
[[0.38026133 0.3381979 0.6418481 0.47090608]]
1/1 [=====] - 0s 21ms/step
[[0.3219837 0.2806773 0.58761704 0.4613152]]
1/1 [=====] - 0s 21ms/step
[[0.52475876 0.4966131 0.80857474 0.62143785]]
1/1 [=====] - 0s 21ms/step
[[0.06559242 0.23240717 0.64683884 0.4655144]]
1/1 [=====] - 0s 22ms/step
[[0.03041572 0.22160062 0.56966877 0.46341726]]
1/1 [=====] - 0s 22ms/step
[[0.46727416 0.47965804 0.7700626 0.5995648]]
1/1 [=====] - 0s 22ms/step
[[0.2003838 0.24612685 0.37321794 0.399646]]
1/1 [=====] - 0s 22ms/step
[[0.3941683 0.4436239 0.64734644 0.5379286]]
1/1 [=====] - 0s 22ms/step
[[0.32886672 0.23317465 0.61476713 0.43255025]]
1/1 [=====] - 0s 22ms/step
[[0.16727684 0.26254746 0.3716781 0.41988584]]

1/1 [=====] - 0s 21ms/step
[[0.5077393 0.48862728 0.79542375 0.60954165]]



In [16]: `model.save('./Saved-Models/Locator')`

INFO:tensorflow:Assets written to: ./Saved-Models/Locator/assets
INFO:tensorflow:Assets written to: ./Saved-Models/Locator/assets